

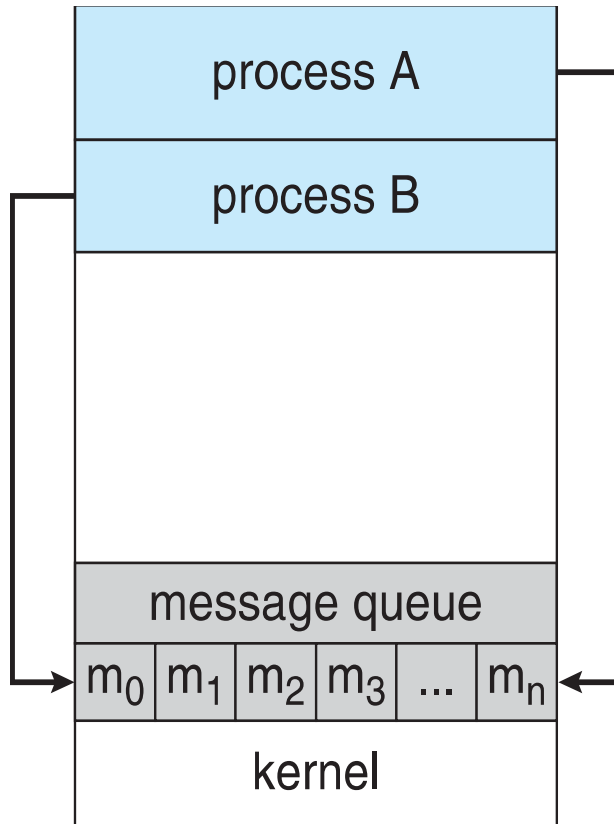
Inter-Process-Communication

Interprocess Communication


- Processes within a system may be *independent* or *cooperating*
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **inter-process communication (IPC)**
- Two models of IPC
 - **Shared memory**
 - **Message passing**

Communications Models

(a) Message passing. (b) shared memory.



(a)



(b)


Shared Memory

One process will create an area in RAM which the other process can access

Both processes can access shared memory like a regular working memory

- Reading/writing is like regular reading/writing
- Fast

Limitation : Error prone. Needs synchronization between processes



Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - **unbounded-buffer** places no practical limit on the size of the buffer
 - **bounded-buffer** assumes that there is a fixed buffer size

Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

Bounded-Buffer – Producer

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Bounded Buffer – Consumer

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next
consumed */
}
```


Shared Memory in Linux

`int shmget (key, size, flags)`

- Create a shared memory segment;
- Returns ID of segment : `shmid`
- `key` : unique identifier of the shared memory segment
- `size` : size of the shared memory (rounded up to the `PAGE_SIZE`)

`int shmat(shmid, addr, flags)`

- **A**ttach `shmid` shared memory to address space of the calling process
- `addr` : pointer to the shared memory address space

`int shmdt(shmid)`

- **D**etach shared memory

server.c

```
1 #include <sys/types.h>
2 #include <sys/ipc.h>
3 #include <sys/shm.h>
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 #define SHMSIZE 27 /* Size of shared memory */
8
9 main()
10 {
11     char c;
12     int shmid;
13     key_t key;
14     char *shm, *s;
15
16     key = 5678; /* some key to uniquely identifies the shared memory */
17
18     /* Create the segment. */
19     if ((shmid = shmget(key, SHMSIZE, IPC_CREAT | 0666)) < 0) {
20         perror("shmget");
21         exit(1);
22     }
23
24     /* Attach the segment to our data space. */
25     if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
26         perror("shmat");
27         exit(1);
28     }
29
30     /* Now put some things into the shared memory */
31     s = shm;
32     for (c = 'a'; c <= 'z'; c++)
33         *s++ = c;
34     *s = 0; /* end with a NULL termination */
35
36     /* Wait until the other process changes the first character
37      * to '*' the shared memory */
38     while (*shm != '*')
39         sleep(1);
40     exit(0);
41 }
```

client.c

```
1 #include <sys/types.h>
2 #include <sys/ipc.h>
3 #include <sys/shm.h>
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 #define SHMSIZE 27
8
9 main()
10 {
11     int shmid;
12     key_t key;
13     char *shm, *s;
14
15     /* We need to get the segment named "5678", created by the server
16     key = 5678;
17
18     /* Locate the segment. */
19     if ((shmid = shmget(key, SHMSIZE, 0666)) < 0) {
20         perror("shmget");
21         exit(1);
22     }
23
24     /* Attach the segment to our data space. */
25     if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
26         perror("shmat");
27         exit(1);
28     }
29
30     /* read what the server put in the memory. */
31     for (s = shm; *s != 0; s++)
32         putchar(*s);
33     putchar('\n');
34
35     /*
36      * Finally, change the first character of the
37      * segment to '*', indicating we have read
38      * the segment.
39      */
40     *shm = '*';
41
42     exit(0);
43 }
```


Message Passing

Shared memory created in the kernel

System calls such as **send** and **receive** used for communication

- Cooperating : each send must have a receive

Advantage : Explicit sharing, less error prone



Message Passing (Cont.)

- Implementation of communication link
 - Physical:
 - Shared memory
 - Hardware bus
 - Network
 - Logical:
 - Direct or indirect
 - Synchronous or asynchronous
 - Automatic or explicit buffering

Direct Communication

- Processes must name each other explicitly:
 - **send** (*P, message*) – send a message to process P
 - **receive**(*Q, message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - A pair of processes can have at most one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Send (A, message) – Send message to mailbox A
 - receive (A, message) – receive message from mailbox A
 - Each mailbox has a unique id
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 , sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Indirect Communication

- Mailbox may be owned either by process or OS.
- If a mailbox is owned by a process then only that process can receive message through that mailbox and it can allow a set of processes to send message via this mailbox.
- With termination of process such mailbox destroyed
- Mailbox owned by OS remains forever.

Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** -- the sender is blocked until the message is received
 - **Blocking receive** -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** -- the sender sends the message and continue
 - **Non-blocking receive** -- the receiver does not do busy waiting, it receives
 - A valid message, or
 - Null message
- Different combinations possible

Buffering

- Queue of messages attached to the link.
- implemented in one of three ways
 1. Zero capacity – no messages are queued on a link. Sender must wait for receiver
 2. Bounded capacity – finite length of n messages. Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits

Pipe


Always **between parent and child**

Always **unidirectional**

Accessed by two associated file descriptors:

- fd[0] for reading from pipe
- fd[1] for writing to the pipe

- ☐ Ordinary Pipes allow communication in standard producer-consumer style
- ☐ Producer writes to one end (the **write-end** of the pipe)
- ☐ Consumer reads from the other end (the **read-end** of the pipe)




Created using *pipe()*:

```
int filedes[1];  
pipe(filedes);
```

...


```
write(filedes[1], buf, count);  
read(filedes[0], buf, count);
```



```
int fildes[2];

pipe(fildes);

child_pid = fork();
if (child_pid == 0) {
    close(fildes[1]);
    /* Child now reads */
} else {
    close(fildes[0]);
    /* Parent now writes */
}
```



Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems