

Chapter 4: Buffer Overflow Attack

Copyright © 2017 Wenliang Du, All rights reserved.

Problems

- 4.1. How are the addresses decided for the following variables?

```
void foo(int a)
{
    int x;
    static int y;
}
```

- 4.2. In which memory segments are the variables in the following code located?

```
int i = 0;
void func(char *str)
{
    char *ptr = malloc(sizeof(int));
    char buf[1024];
    int j;
}
```

- 4.3. Please draw the function stack frame for the following C function.

```
int bof(char *str)
{
    char buffer[24];
    strcpy(buffer, str);
    return 1;
}
```

- 4.4. A student proposes to change how the stack grows. Instead of growing from high address to low address, the student proposes to let the stack grow from low address to high address. This way, the buffer will be allocated above the return address, so overflowing the buffer will not be able to affect the return address. Please comment on this proposal.
- 4.5. In the buffer overflow example shown in Listing 4.1, the buffer overflow occurs inside the `strcpy()` function, so the jumping to the malicious code occurs when `strcpy()` returns, not when `foo()` returns. Is this true or false? Please explain.
- 4.6. The buffer overflow example was fixed as below. Is this safe ?

```
int bof(char *str, int size)
{
    char *buffer = (char *) malloc(size);

    /* The following statement has a buffer overflow problem */
```

```

    strcpy(buffer, str);

    return 1;
}

```

- 4.7. In `exploit.c` (Listing 4.2), when assigning the value for the return address, can we do the following? Do you think the return address will point to the shell code or not? Why?

```

*((long *) (buffer + 0x24)) = buffer+ 0x150;

```

- 4.8. Several students had issue with the buffer overflow attack. Their badfile was constructed properly where shell code is at the end of badfile, but when they try different return addresses, they get the following observations. Can you explain why some addresses work and some do not?

```

buffer address : 0xbffff180
case 1 : long retAddr = 0xbffff250 -> Able to get shell access
case 2 : long retAddr = 0xbffff280 -> Able to get shell access
case 3 : long retAddr = 0xbffff300 -> Cannot get shell access
case 4 : long retAddr = 0xbffff310 -> Able to get shell access
case 5:  long retAddr = 0xbffff400 -> Cannot get shell access

```

- 4.9. The following function is called in a privileged program. The argument `str` points to a string that is entirely provided by users (the size of the string is up to 300 bytes). When this function is invoked, the address of the `buffer` array is `0xAABB0010`, while the return address is stored in `0xAABB0050`. Please write down the string that you would feed into the program, so when this string is copied to `buffer` and when the `bof()` function returns, the privileged program will run your code. In your answer, you don't need to write down the injected code, but the offsets of the key elements in your string need to be correct. Note: there is a trap in this problem; some people may be lucky and step over it, but some people may fall into it. Be careful.

```

int bof(char *str)
{
    char buffer[24];
    strcpy(buffer, str);
    return 1;
}

```

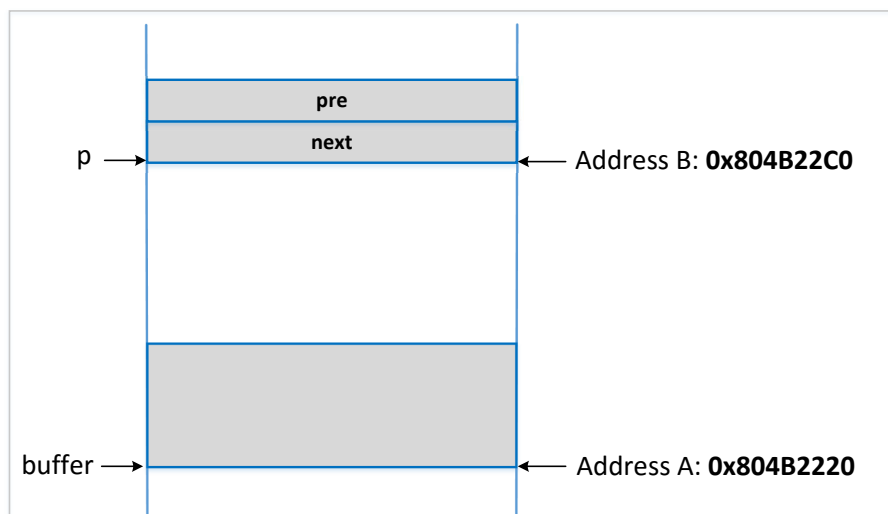


Figure 1: Figure for Problem 4.10.

4.10. ★ ★ ★

In this problem, we will figure out how overflowing a buffer on the heap can lead to the execution of malicious code. The following is a snippet of a code execution sequence (not all code in this sequence is shown here). During the execution of this sequence, the memory locations of `buffer` and the node `p`, which are allocated on the heap, are depicted in Figure 1. You can provide your input (up to 300 bytes) in `user_input`, which will be copied to `buffer`. Your job is to overflow the buffer, so when the target program gets to Line ❸, it will jump to the code that you have injected into the heap memory (assuming that the heap memory is executable). The return address is stored at location `0xBBFFFAACC`.

```
struct Node
{
    struct Node *next;
    struct Node *pre;
};

// The following is a snippet of a code execution sequence.

struct Node *p = malloc(sizeof(struct Node));
struct Node *q;
char *buffer = malloc(100);

/* Code omitted: Add Node p to a linked list */

// There is a potential buffer overflow in the following
strcpy(buffer, user_input);

// remove Node p from the linked list
```

```

q = p->pre;           ❶
q->next = p->next;     ❷

return;               ❸

```

Hint: You still want to place the starting address of your malicious code into the return address field located at `0xBBFFAACC`. Unlike stack-back buffer overflows, where you can naturally reach the return address field via overflowing, now the buffer is on the heap, but the return address is on the stack; you cannot reach the stack by overflowing something on the heap. You should take advantage of the operations on the linked list (Lines ❶ and ❷) to modify the return address field.

This is a simplified version of how a buffer overflow on the heap can be exploited. The linked list is not part of the vulnerable program; it is actually part of the operating system, which uses it to manage the memory on the heap for the current process. Unfortunately, the linked list is also stored on the heap, so by overflowing an application's buffer, attackers can change the values on this linked list. When the OS operates on the corrupted linked list, it may change the return address of function, and trigger the execution of the injected code.

- 4.11. This problem is built on top of Problem 4.10.. Assume that the structure for `Node` becomes the following (a new integer field is added to the beginning). Please redo Problem 4.10.. It should be noted that in Figure 1, the variable `p` will now point to the area 4 bytes below the `next` field.

```

struct Node
{
    int value;
    struct Node *next;
    struct Node *pre;
};

```

- 4.12. Why does ASLR make buffer-overflow attack more difficult?

- 4.13. ★

To write a shellcode, we need to know the address of the string `"/bin/sh"`. If we have to hardcode the address in the code, it will become difficult if ASLR is turned on. Shellcode solved that problem without hardcoding the address of the string in the code. Please explain how the shellcode in `exploit.c` (Listing 4.2) achieved that.